



G11: GRPO only on prompts G8 solves sometimes (87% generator) vs G10: all prompts (71% Dolci)

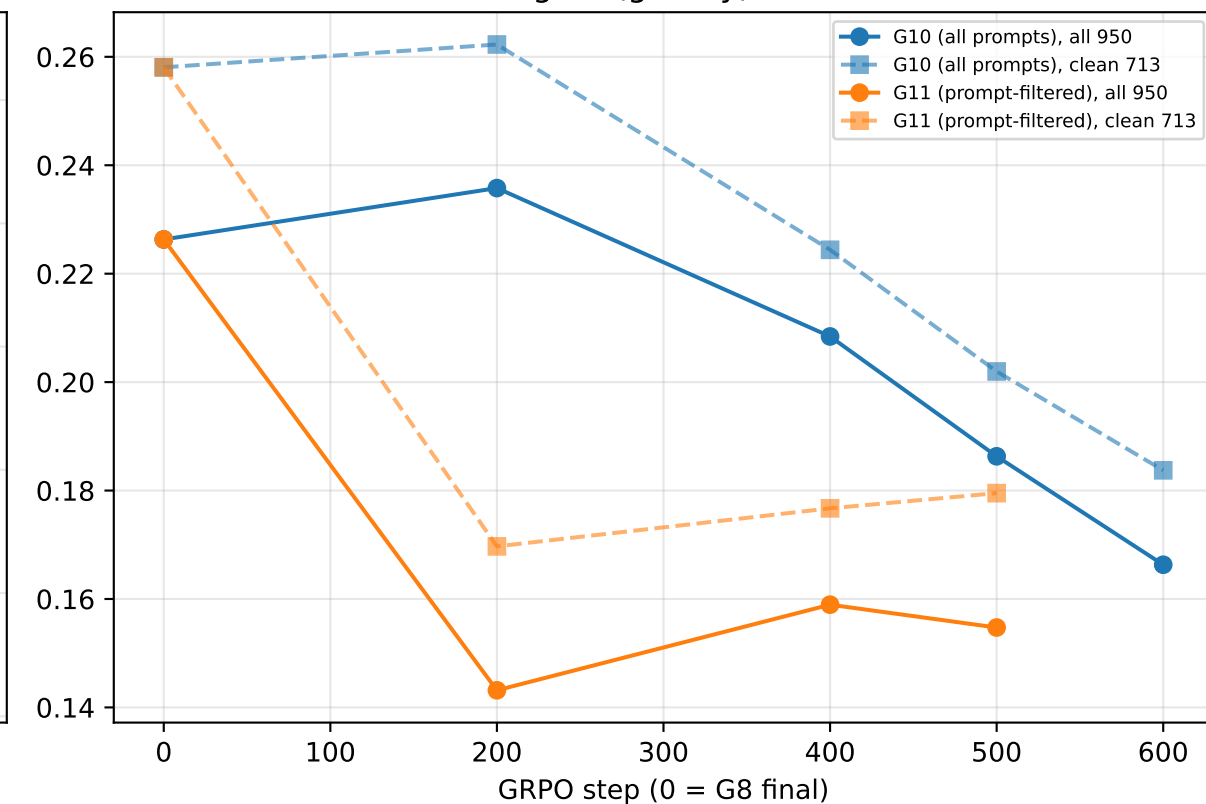
Dolci gate (greedy): valid



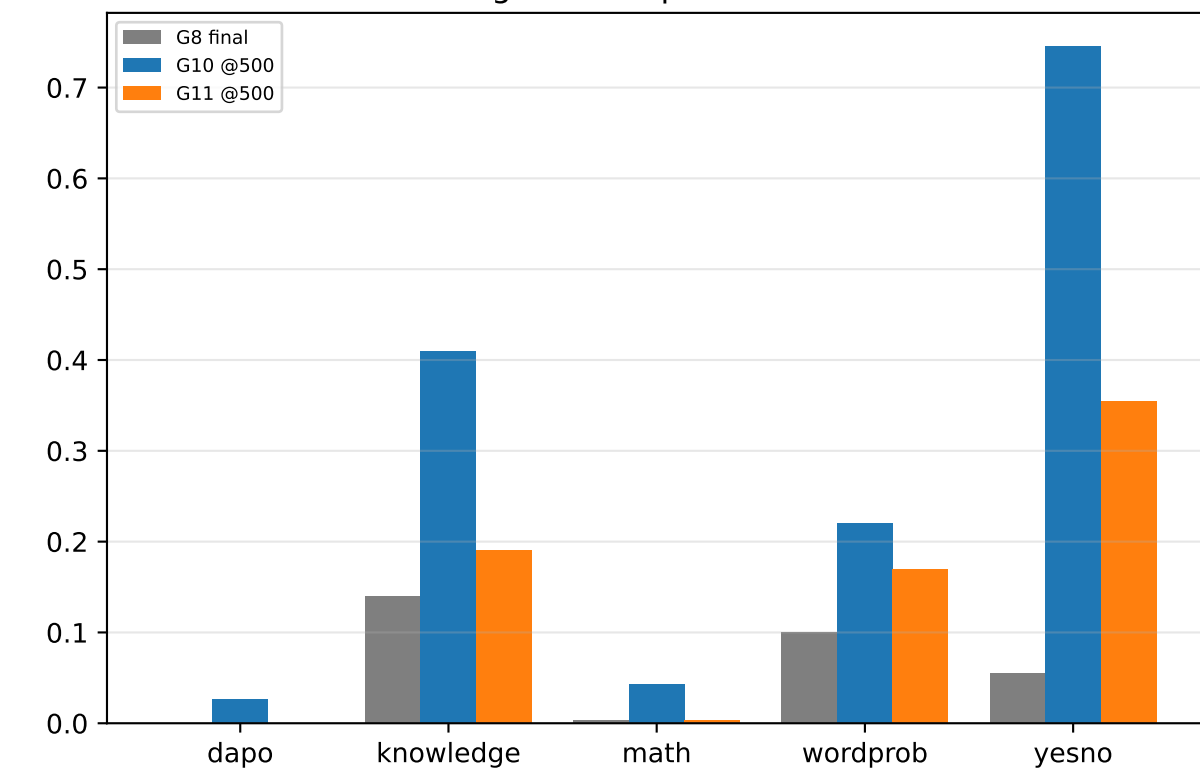
Dolci gate (greedy): valid_correct



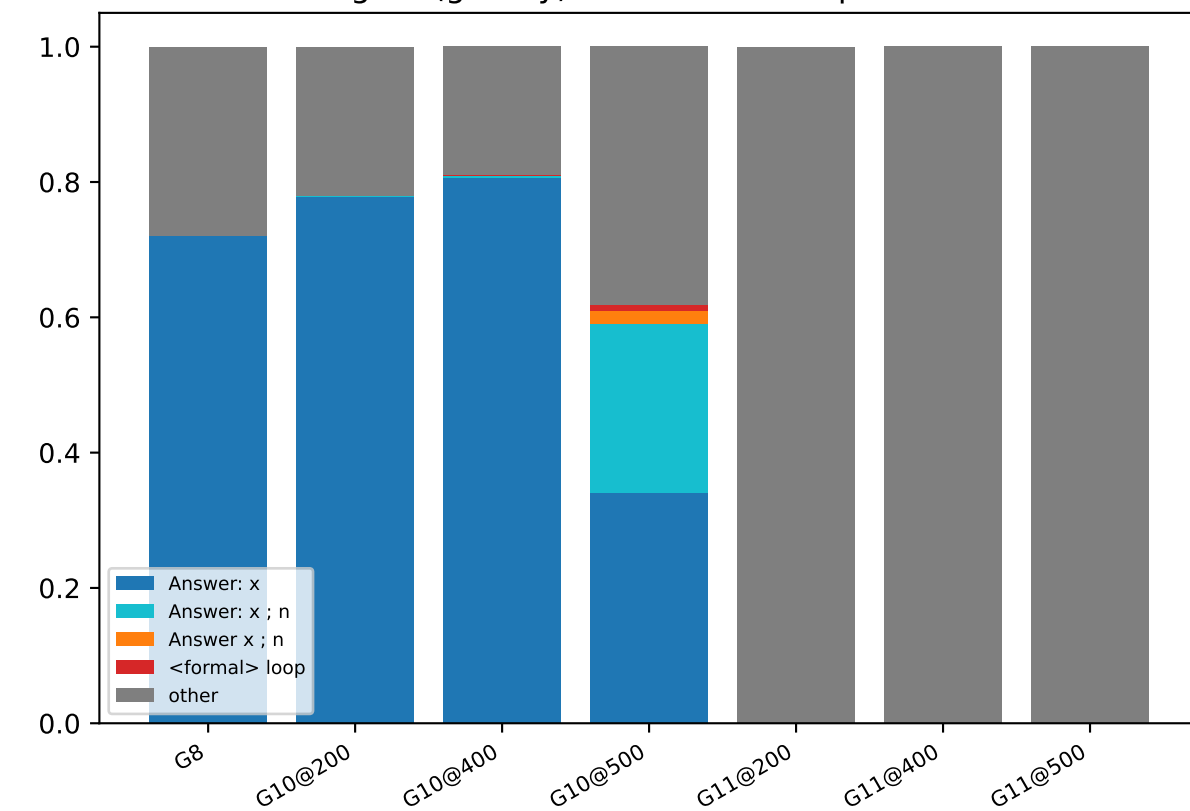
Dolci gate (greedy): correct



gate valid per bench



gate (greedy): what follows </proof>



training rollouts (own prompt pools, not comparable in level)

