



# GRPO with completions cut at 2048 tokens masked from the loss (TRL default): truncation, entropy, loops

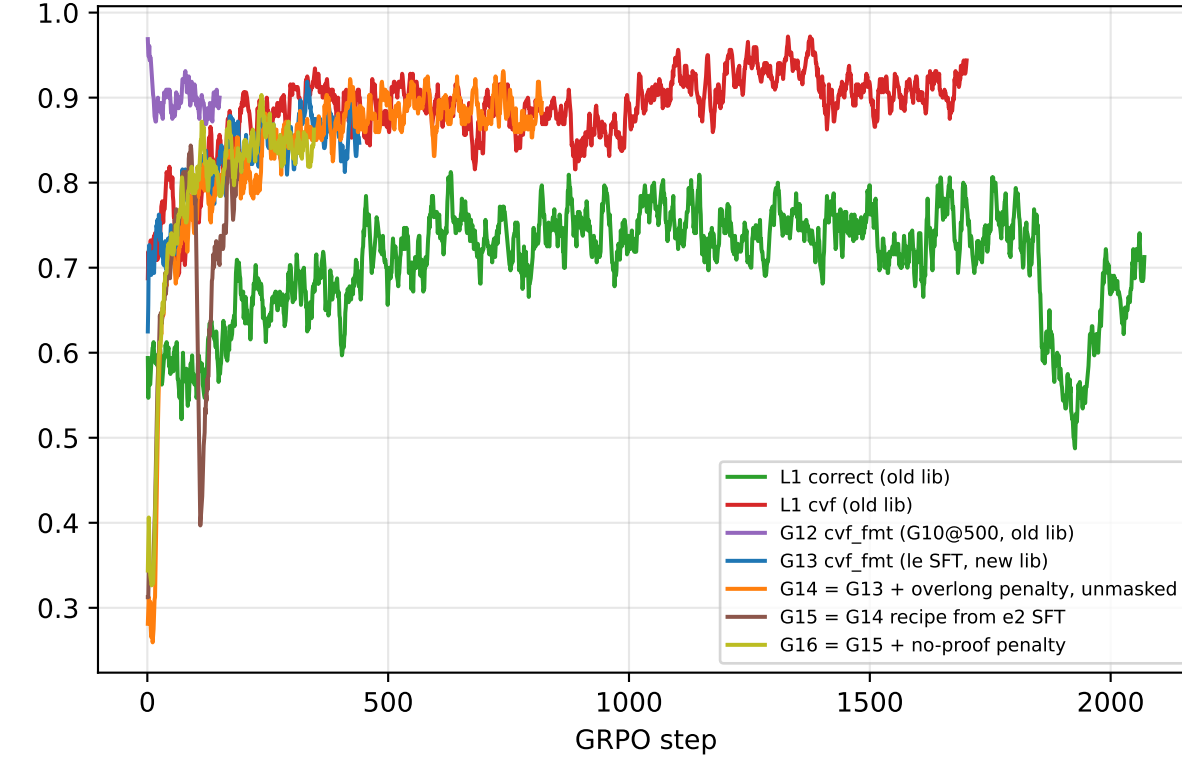
## truncated at 2048 tokens (clipped ratio), 10-step mean



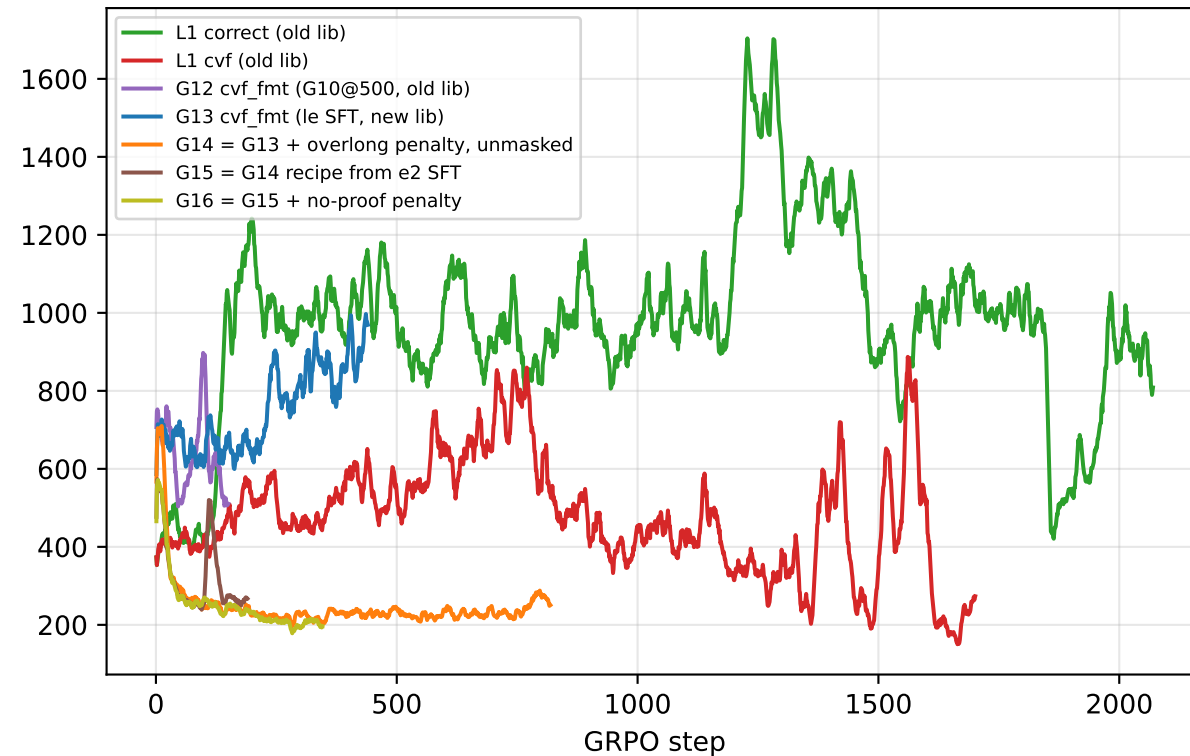
## policy entropy, 10-step mean



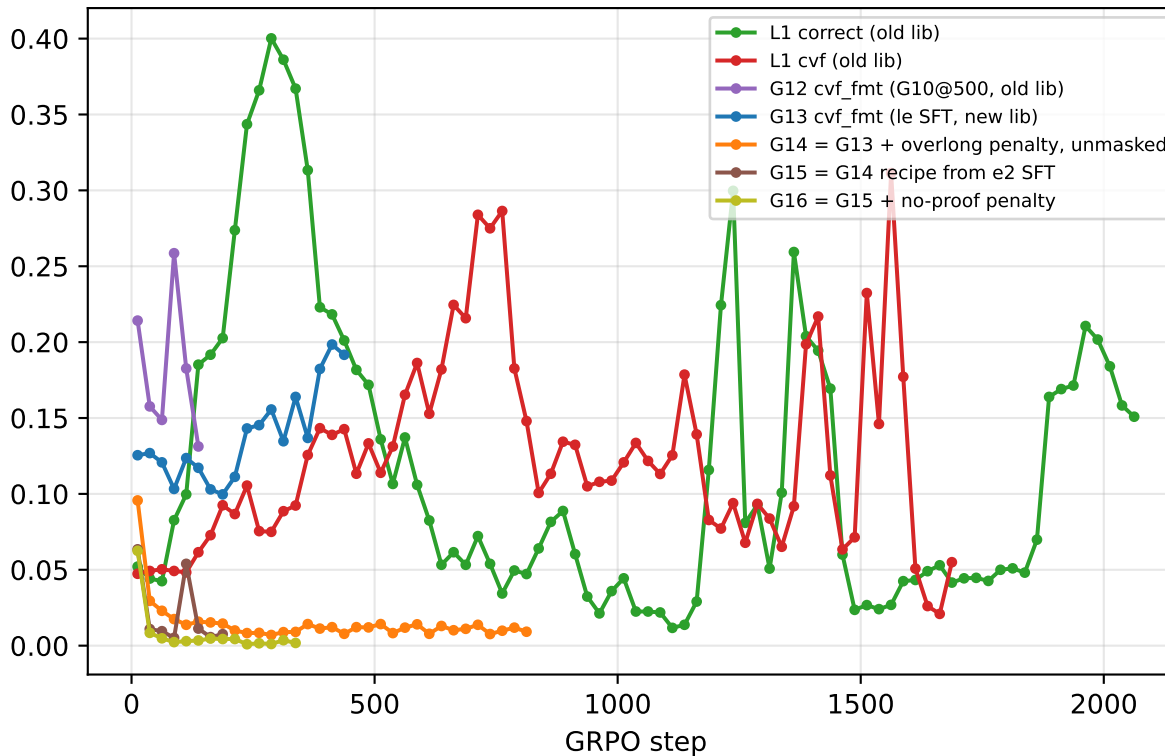
## groups with zero reward variance, 10-step mean



## mean completion length (tokens), 10-step mean



## training rollouts with a line repeated $\geq 8x$ (digits masked)



## training-rollout reward component (own prompt pools)

