



GRPO with completions cut at 2048 tokens masked from the loss (TRL default): truncation, entropy, loops


truncated at 2048 tokens (clipped ratio), 10-step mean




policy entropy, 10-step mean




groups with zero reward variance, 10-step mean



mean completion length (tokens), 10-step mean



training rollouts with a line repeated ≥ 8 x (digits masked)



training-rollout reward component (own prompt pools)

